Department of Digital Systems
MSC in Artificial Intelligence

Fake News Prediction

Dimitrios Delikonstantis

14/02/2021

Reason for subject selection

- Fake or misleading news can be dangerous
- It is also used for making money via advertising (clickbait)
- Prevalence has increased with the rise of social media
- Social media algorithms have been implicated in the spread of fake news
- Since coronavirus appeared fake news have been on the rise
- Anti-vaccination movement is growing vastly via fake news, meanwhile WHO declared it one of the top ten global threats
- Therefore, vastly reducing fake news is crucial for the society's own good

- Dataset retrieved from kaggle
- Used data from 20800 article news labeled 'Real' or 'Fake'

Bag-of-words model

- Text is changed into vectors of numbers to be processed by machine learning algorithms (feature encoding)

- Describes the occurrence of words within a document

- Transform each document from a corpus of documents into a vector and use it as an input to a machine learning algorithm

Feature extraction with TF-IDF

- Statistical measure that shows how important a word is to a document in an entire corpus of documents
- Term frequency shows the frequency of a word in a document
- TF = Word occurrence / vocabulary of document
- Inverse document frequency shows the frequency of a word within the corpus of documents.
- IDF = log(Entire corpus of documents / number of documents containing specific word)
- TF-IDF score = TF * IDF
- The higher the score, the more relevant that word is in that particular document
- Sklearn TfidfVectorizer used for different case scenarios

Train test split and cross-validation

- Sklearn train_test_split was used for splitting the dataset. Different case scenarios were tested to see which one provided best results.
- But what if the training dataset includes only articles from a specific author. Then our data is biased
- This is why cross-validation is crucial
- Cross-validation (partially) was implemented with sklearn Kfold to check the accuracy on the training data and compare it to the predicted data to check if overfitting occurs
- Cross-validation will be extended to the code to train and evaluate the models used and see which algorithm performs better for our application
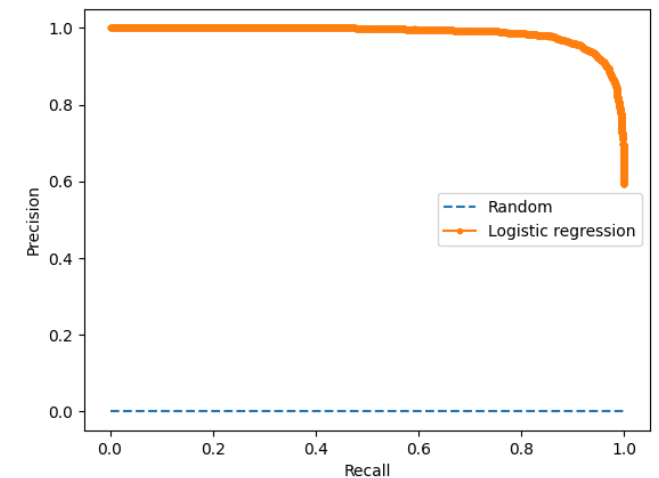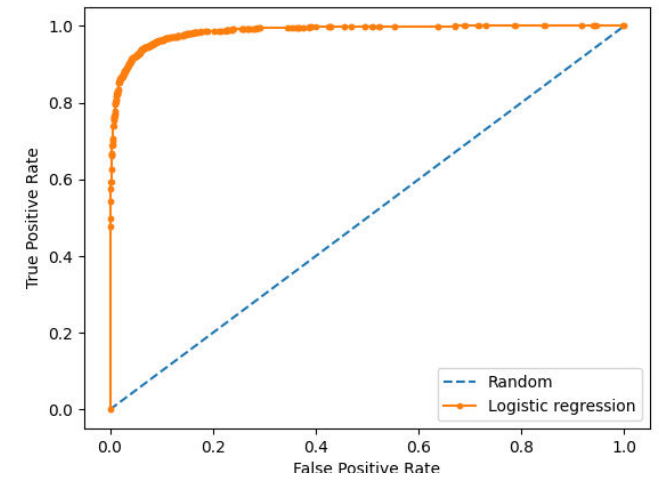
Logistic regression metrics

- Cross-validation mean accuracy: 93.2%
- Accuracy: 93.55%
- Confusion matrix:

		Actual class
Predicted		[[1982 118]
class		[150 1910]]

- Precision: 0.9418
- Recall: 0.9271
- Specificity: 0.9438
- False positive rate: 0.056
- F1-score: 0.9344

- ## ROC AUC: 0.9853

- ## ROC AUC: 0.9857

Naïve Bayes metrics

- Cross-validation mean accuracy: 87.76%
- Accuracy: 87.21%
- Confusion matrix:

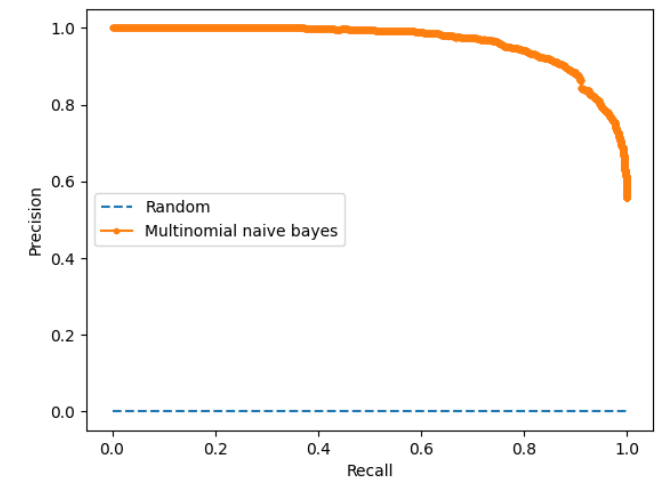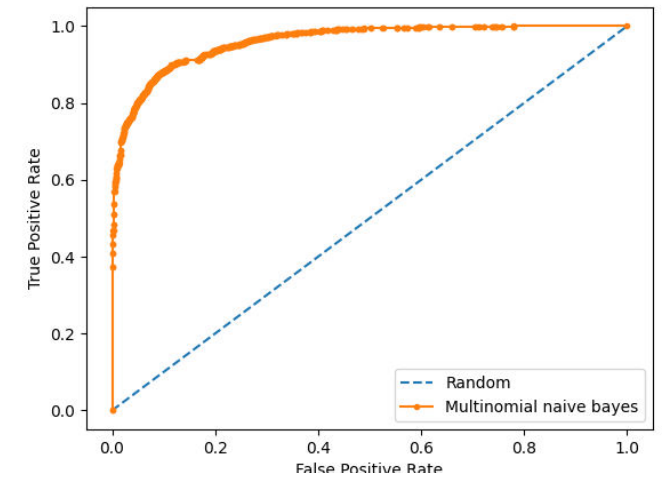|	Actual class
Predicted class	[[1752 348] [184 1876]]

- Precision: 0.8435
- Recall: 0.9106
- Specificity: 0.8342
- False positive rate: 0.1657
- F1-score: 0.8758

- ROC AUC: 0.9601

- ROC AUC: 0.9626

K-nearest neighbor metrics

- Cross-validation mean accuracy: 72.51%
- Accuracy: 73.72%
- Confusion matrix:

	Actual class
Predicted	[[1994 106]
class	[987 1073]]

- Precision: 0.91
- Recall: 0.5208
- Specificity: 0.9495
- False positive rate: 0.05
- F1-score: 0.6625

K-nearest neighbor ROC/Precision recall curves

- ROC AUC: 0.7953

- ROC AUC: 0.8511
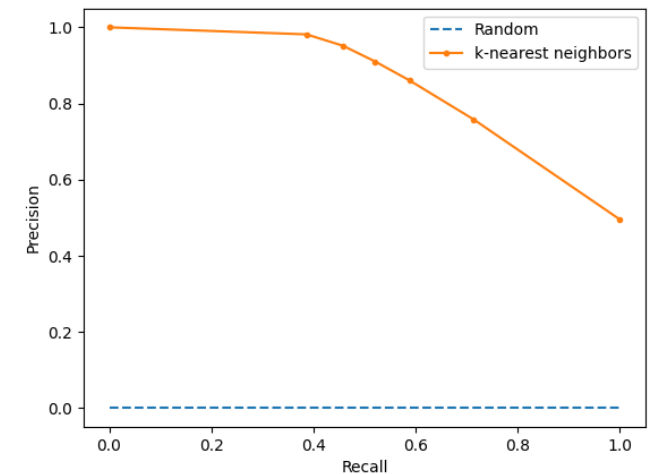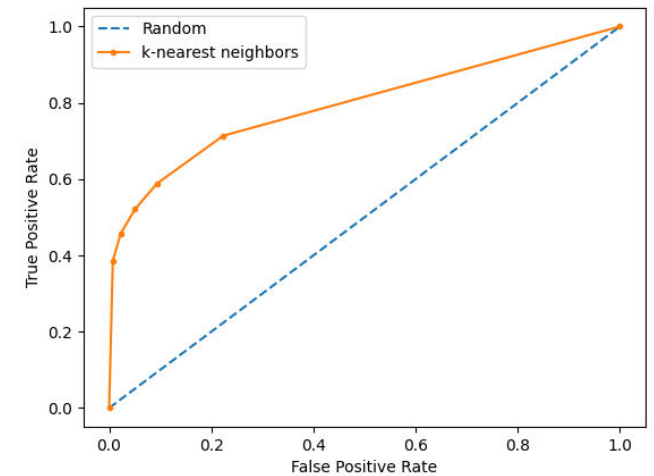
Decision tree metrics

- Cross-validation mean accuracy: 87.6%
- Accuracy: 87.3%
- Confusion matrix:

		Actual class
Predicted class		[[1822 278]
		[250 1810]]

- Precision: 0.8668
- Recall: 0.8786
- Specificity: 0.8676
- False positive rate: 0.1323
- F1-score: 0.8727

- # ROC AUC: 0.8714

- # ROC AUC: 0.9012